

Bijon Setyawan Raya

bijonsetyawan@gmail.com | github.com/bsraya | <https://bijonsetyawan.com>

WORK EXPERIENCE

Senior Machine Learning Engineer

Oct. 2025 — Present

IKG Limited Group

Taipei, Taiwan

- Designed and shipped a content moderation system from scratch, replacing an LLM-based PoC with a Go-based embedding + reranker pipeline — reducing response time by 99% (8s → 785ms) and reducing inference costs by 99% (100k to <1k tokens)
- Built a precomputed financial metrics layer using tiered data architecture (Bronze/Silver/Gold), reducing 70% avg. query time (10s+ to <3s)
- Architected and deployed the first automated data pipeline on all environments, transforming raw source data into governed financial metrics at 1M+ records/day — maintaining 99.9% SLA

Machine Learning Engineer

Nov. 2024 — Oct. 2025

Cathay United Bank

Taipei, Taiwan

- Architected a 3-tier content safety system, reducing response time from 50s+ to under 5s while blocking 99% of unsafe interactions
- Prototyped a no-code MLOps pipeline enabling scientists to train, validate, and deploy scikit-learn models, cut training time 5× with cuML and automated model deployment for real-time inference

Fullstack Developer

Oct. 2022 — Mar. 2024

Faria Education Group

Taipei, Taiwan

- Migrated 25% of legacy CoffeeScript codebase to Stimulus.js, cutting build times by 40%
- Optimized file export pipeline, reducing export times from 5+ minutes to under 3 seconds (99% faster)

SKILLS

Languages: Go, Python, SQL

Deep Learning & Machine Learning: CUDA, HIP/ROCm, MLflow, cuML, PyTorch, Scikit-learn

Data Engineering: Airflow, Iceberg, dbt, DuckDB, MongoDB, PostgreSQL, Redis, Qdrant

System & Infrastructure: Docker, Git, Kubernetes, Terraform, Linux/Unix, Podman

Cloud Services: Amazon Web Service, Google Cloud Provider

PROJECTS

Chat Moderation System & Guardrails

- An AI-powered detection system filtering harmful messages in real-time across platforms
- A prompt filtering system that ensures safe LLM interactions

Distributed Machine Learning System

- A MLOps platform that simplifies ML model life cycles

Vanilla RAG & Multi-Modal RAG

- High-fidelity retrieval RAG systems for complex data

Image Search Engine

- A search engine with a 95% reduction in search latency and 87% lower storage footprint

CERTIFICATES

Machine Learning Specialty - Amazon Web Service

Oct. 2025

EDUCATION

National Tsing Hua University

Hsinchu, Taiwan

Master of Science in Computer Science and Mathematics

Feb. 2021 — Jan. 2023

National Tsing Hua University

Hsinchu, Taiwan

Bachelor of Science in Computer Science and Mathematics

Sep. 2015 — Jan. 2021